

Lab7

Jun Rao

10/8/2019

Hierarchical linear regression

In this lab we will test our understanding of random effects (mixed-effects) modeling.

```
#-----  
#-----  
# Import/Load the rstan library:  
library(tidyverse)  
library(rstan)  
library(loo)  
rstan_options(auto_write = TRUE)  
options(mc.cores = parallel::detectCores())  
#-----  
#-----  
## Simulating the data
```

Simulating the data

Our first task is to simulate data following the assumptions of random-effects modeling. First, we will assume a true model in which groups have an equal slope, but vary in their intercept:

$$y_i = (\bar{\alpha} + \eta_{j,[i]}) + \beta x_i$$

Task 1 (25 points)

Complete the following:

1. Add comments to this code-chunk to convey your understanding (2 points).

```
#creating simulations or random objects that can be reproduced.  
set.seed(7)
```

```
#We have 8 groups  
n_group = 8
```

```
#how many data points in each group  
n_samp_per_group = ceiling(runif(n_group, 4, 15))
```

```
group_idx = rep(c(1:n_group), times = n_samp_per_group)
```

```
#total data points in all groups  
n_sample = sum(n_samp_per_group)
```

```
#Each Parameter in the model need a prior  
alpha_mean = rnorm(1, 0, 10)  
alpha_sigma = abs(rnorm(1, 0, 10))  
alpha_eta_vec = rnorm(n_group, 0, alpha_sigma)
```

```
beta = rnorm(1, 0, 10)
```

```
sigma_resid = abs(rnorm(1, 0, 2))
```

```
#input value  
x_raw = runif(n_sample, 0, 3500)  
x_mean = mean(x_raw)  
x_sd = sd(x_raw)
```

```
#center and standardize:  
x_scaled = (x_raw - x_mean) / (2 * x_sd)
```


```
epsilon = rnorm(n_sample, 0, sigma_resid)
```

```
y_hat = NULL  
for(i in 1:n_sample){  
  y_hat[i] = (alpha_mean + alpha_eta_vec[group_idx[i]]) + x_scaled[i] * beta  
}
```

```
#Output Value  
y_obs = y_hat + epsilon
```

```
data_df =
  data.frame(y_obs, x_scaled, group_idx) %>%
  mutate(group_idx = factor(group_idx))

#draw the graphics for each group
ggplot(data_df, aes(x = x_scaled, y = y_obs, color = group_idx)) +
  geom_point(shape = 19) +
  geom_smooth(method = "lm", se = FALSE, size = 0.2) +
  scale_color_brewer(palette = "Set1") +
  geom_vline(xintercept = 0, linetype = 2) +
  theme_classic()
```



Fitting the Stan models

2. Use the supplied Stan code to fit the random effects model to your generated data.:

Remember to monitor your variable so that you can complete sub-task 5, below (5 points).

```
stan_data = list(n_sample = n_sample,
  n_group = n_group,
  y_vec = y_obs,
  x_vec = x_scaled,
  group_idx = group_idx)
```

```
params_monitor = c("beta", "alpha_mean", "eta_alpha", "alpha_sigma", "sigma_r",
  "esid", "log_lik")
```

```

test_fit_1 = stan( file = "Intercepts_LinReg.stan",
                  data = stan_data,
                  pars = params_monitor,
                  chains = 1,
                  iter = 10,
                  algorithm="NUTS")

##
## SAMPLING FOR MODEL 'Intercepts_LinReg' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: WARNING: No variance estimation is
## Chain 1:           performed for num_warmup < 20
## Chain 1:
## Chain 1: Iteration: 1 / 10 [ 10%] (Warmup)
## Chain 1: Iteration: 2 / 10 [ 20%] (Warmup)
## Chain 1: Iteration: 3 / 10 [ 30%] (Warmup)
## Chain 1: Iteration: 4 / 10 [ 40%] (Warmup)
## Chain 1: Iteration: 5 / 10 [ 50%] (Warmup)
## Chain 1: Iteration: 6 / 10 [ 60%] (Sampling)
## Chain 1: Iteration: 7 / 10 [ 70%] (Sampling)
## Chain 1: Iteration: 8 / 10 [ 80%] (Sampling)
## Chain 1: Iteration: 9 / 10 [ 90%] (Sampling)
## Chain 1: Iteration: 10 / 10 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.001 seconds (Warm-up)
## Chain 1:           0 seconds (Sampling)
## Chain 1:           0.001 seconds (Total)
## Chain 1:

## Now we will run our full model:
# How many samples do we want of each parameter, from each chain?
n_mc_samples = 1000
# How much burn-in?
n_burn = 500
# How much thinning? (take the ith value of the chain)
n_thin = 3
# Total iterations needed:
n_iter_total = (n_mc_samples * n_thin) + n_burn
model_fit_1 =
stan(fit = test_fit_1, # So it knows we're already compiled
    file = "Intercepts_LinReg.stan",
    data = stan_data,
    pars = params_monitor,

```

```

    chains = 3,
    warmup = n_burn,
    thin = n_thin,
    iter = n_iter_total,
    algorithm="NUTS")
model_out_1 = rstan::extract(model_fit_1)
str(model_out_1)

## List of 7
## $ beta      : num [1:3000(1d)] 19.7 19.4 19.4 19.3 19.4 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ alpha_mean : num [1:3000(1d)] 0.329 6.789 3.795 -6.03 -0.922 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ eta_alpha  : num [1:3000, 1:8] -2.37 -8.9 -6.18 3.61 -1.59 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ iterations: NULL
## .. ..$      : NULL
## $ alpha_sigma: num [1:3000(1d)] 14.3 8.43 9.38 12.24 8.07 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ sigma_resid: num [1:3000(1d)] 0.727 0.761 0.792 0.95 0.837 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ log_lik     : num [1:3000, 1:78] -0.723 -0.678 -0.906 -0.98 -1.024 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ iterations: NULL
## .. ..$      : NULL
## $ lp__        : num [1:3000(1d)] -50.1 -45.6 -47 -50.3 -45.5 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL

```

3. Construct at least two visualizations to validate that your model was able to adequately recover the true parameters. Note: A table can be considered a visualization, because it organizes the quantitative information in a visually accessible manner (5 points).

#first we visualize tabularly, the 95% CI for estimated parameters and compare these to the true values

```

model_sum = summary(model_fit_1)$summary
credint = model_sum[c(1:12),c(4,6,8)]
credint2 = cbind(credint, c(beta, alpha_mean, alpha_sigma, alpha_eta_vec[1:8],
sigma_resid))
colnames(credint2)[4] <- "true_vals"

```

correct true eta values

```

corrected = c(beta, alpha_mean + mean(alpha_eta_vec), alpha_sigma, alpha_eta_vec + alpha_mean, sigma_resid)
credint2 = cbind(credint2, corrected)

```

```
colnames(credint2)[5] <- "corrected"
credint2
```

	2.5%	50%	97.5%	true_vals	corrected
## beta	18.9236199	19.3179289	19.7174351	18.960671	18.9606707
## alpha_mean	-7.5914333	0.5827238	8.5616500	-9.706733	0.5393192
## eta_alpha[1]	-10.8982610	-2.9252298	5.2614329	9.472799	9.4727995
## eta_alpha[2]	-20.2605034	-12.2285380	-4.1027645	7.086974	-2.6197595
## eta_alpha[3]	-16.4345432	-8.4601092	-0.3469185	-1.107893	-10.8146268
## eta_alpha[4]	1.9911077	9.9272399	18.0476215	1.446095	-8.2606383
## eta_alpha[5]	-14.4028674	-6.4648197	1.7754864	20.745223	11.0384900
## eta_alpha[6]	7.3278001	15.3823418	23.5730000	3.381659	-6.3250744
## eta_alpha[7]	3.5685271	11.5971335	19.7853115	25.735245	16.0285114
## eta_alpha[8]	-14.9054012	-6.8510395	1.2482065	21.611737	11.9050031
## alpha_sigma	6.7007690	10.4672809	18.5512442	3.069382	-6.6373518
## sigma_resid	0.6984777	0.8212417	0.9790730	0.935361	0.9353610

we can see that few parameters were adequately recovered. This included alpha_mean and all of the eta_alphas (9 of 12 params)

```
par(mfrow = c(5,3))
hist(model_out_1$beta,main="the marginal posterior samples of beta",xlab = expression(beta))
abline(v = beta, lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$alpha_mean,main="the marginal posterior samples of alpha_mean",xlab = expression(alpha_mean))
abline(v = alpha_mean, lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,1],main="the marginal posterior samples of eta_alpha[1]",xlab = expression(eta_alpha[1]))
abline(v = alpha_eta_vec[1], lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,2],main="the marginal posterior samples of eta_alpha[2]",xlab = expression(eta_alpha[2]))
abline(v = alpha_eta_vec[2], lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,3],main="the marginal posterior samples of eta_alpha[3]",xlab = expression(eta_alpha[3]))
abline(v = alpha_eta_vec[3], lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,4],main="the marginal posterior samples of eta_alpha[4]",xlab = expression(eta_alpha[4]))
abline(v = alpha_eta_vec[4], lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,5],main="the marginal posterior samples of eta_alpha[5]",xlab = expression(eta_alpha[5]))
abline(v = alpha_eta_vec[5], lty = 1, lwd = 2, col = "blue")
```

```
hist(model_out_1$eta_alpha[,6],main="the marginal posterior samples of eta_alpha[6]",xlab = expression(eta_alpha[6]))
```

```

pha[6]",xlab = expression(eta_alpha[6]))
abline(v = alpha_eta_vec[6], lty = 1, lwd = 2, col = "blue")

hist(model_out_1$eta_alpha[,7],main="the marginal posterior samples of eta_alpha[7]",xlab = expression(eta_alpha[7]))
abline(v = alpha_eta_vec[7], lty = 1, lwd = 2, col = "blue")

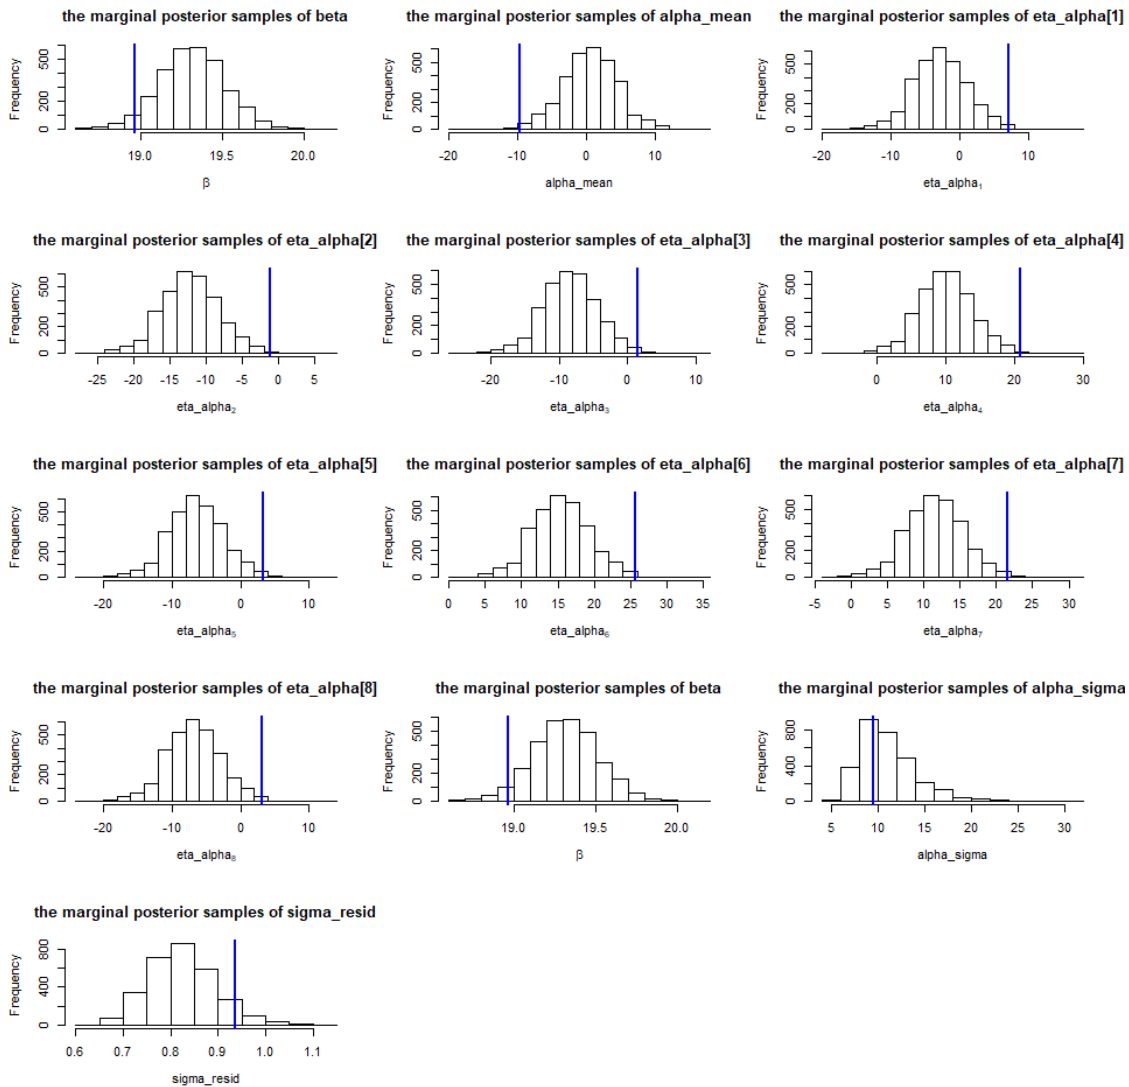
hist(model_out_1$eta_alpha[,8],main="the marginal posterior samples of eta_alpha[8]",xlab = expression(eta_alpha[8]))
abline(v = alpha_eta_vec[8], lty = 1, lwd = 2, col = "blue")

hist(model_out_1$beta,main="the marginal posterior samples of beta",xlab = expression(beta))
abline(v = beta, lty = 1, lwd = 2, col = "blue")

hist(model_out_1$alpha_sigma,main="the marginal posterior samples of alpha_sigma",xlab = expression(alpha_sigma))
abline(v = alpha_sigma, lty = 1, lwd = 2, col = "blue")

hist(model_out_1$sigma_resid,main="the marginal posterior samples of sigma_resid",xlab = expression(sigma_resid))
abline(v = sigma_resid, lty = 1, lwd = 2, col = "blue")

```



4. Construct a Stan code file that allows for both random intercepts and random slopes. Fit this new, more complex model to your same data set (8 points).

```
beta_mean = rnorm(1, 0, 10)
beta_sigma = abs(rnorm(1, 0, 10))
beta_eta_vec = rnorm(n_group, 0, alpha_sigma)

params_monitor = c("beta_mean", "alpha_mean", "eta_beta", "eta_alpha", "beta_sigma", "alpha_sigma", "sigma_resid", "log_lik")

test_fit_2 = stan( file = "Lab7.stan",
  data = stan_data,
  pars = params_monitor,
  chains = 1,
  iter = 10,
  algorithm="NUTS")
```



```

##
## SAMPLING FOR MODEL 'Lab7' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: WARNING: No variance estimation is
## Chain 1: performed for num_warmup < 20
## Chain 1:
## Chain 1: Iteration: 1 / 10 [ 10%] (Warmup)
## Chain 1: Iteration: 2 / 10 [ 20%] (Warmup)
## Chain 1: Iteration: 3 / 10 [ 30%] (Warmup)
## Chain 1: Iteration: 4 / 10 [ 40%] (Warmup)
## Chain 1: Iteration: 5 / 10 [ 50%] (Warmup)
## Chain 1: Iteration: 6 / 10 [ 60%] (Sampling)
## Chain 1: Iteration: 7 / 10 [ 70%] (Sampling)
## Chain 1: Iteration: 8 / 10 [ 80%] (Sampling)
## Chain 1: Iteration: 9 / 10 [ 90%] (Sampling)
## Chain 1: Iteration: 10 / 10 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.001 seconds (Warm-up)
## Chain 1: 0.002 seconds (Sampling)
## Chain 1: 0.003 seconds (Total)
## Chain 1:

## Now we will run our full model:
# How many samples do we want of each parameter, from each chain?
n_mc_samples = 1000
# How much burn-in?
n_burn = 500
# How much thinning? (take the ith value of the chain)
n_thin = 3
# Total iterations needed:
n_iter_total = (n_mc_samples * n_thin) + n_burn
model_fit_2 =
stan(fit = test_fit_2, # So it knows we're already compiled
    file = "Intercepts_LinReg.stan",
    data = stan_data,
    pars = params_monitor,
    chains = 3,
    warmup = n_burn,
    thin = n_thin,
    iter = n_iter_total,
    algorithm="NUTS")
model_out_2 = rstan::extract(model_fit_2)
str(model_out_2)

```

```
## List of 9
## $ beta_mean : num [1:3000(1d)] 19.1 19.2 19.5 19.5 19 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ alpha_mean : num [1:3000(1d)] -0.16 2.51 0.477 9.372 -0.17 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ eta_beta : num [1:3000, 1:8] 0.114 -0.32 -0.286 -0.198 -0.056 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ iterations: NULL
## .. ..$ : NULL
## $ eta_alpha : num [1:3000, 1:8] -1.91 -4.58 -2.35 -11.76 -2.51 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ iterations: NULL
## .. ..$ : NULL
## $ beta_sigma : num [1:3000(1d)] 0.623 0.211 0.113 0.443 0.245 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ alpha_sigma: num [1:3000(1d)] 9.82 6.73 9.3 10.84 16.73 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ sigma_resid: num [1:3000(1d)] 0.746 0.778 0.821 0.676 0.881 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
## $ log_lik : num [1:3000, 1:78] -0.627 -0.682 -0.731 -0.76 -0.967 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ iterations: NULL
## .. ..$ : NULL
## $ lp_ : num [1:3000(1d)] -48 -36.6 -38 -44.2 -48.1 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ iterations: NULL
```

```
#summary(model_fit_2)$summary
```

```
#first we visualize tabularly, the 95% CI for estimated parameters and compare these the true values
```

```
model_sum = summary(model_fit_2)$summary
credint = model_sum[c(1:21),c(4,6,8)]
credint2 = cbind(credint, c(beta_mean, alpha_mean, beta_sigma, alpha_sigma, beta_eta_vec[1:8], alpha_eta_vec[1:8], sigma_resid))
colnames(credint2)[4] <- "true_vals"
```

```
# correct true eta values
```

```
corrected = c(beta_mean+mean(beta_eta_vec), alpha_mean + mean(alpha_eta_vec),
alpha_sigma, beta_eta_vec + beta_mean, alpha_eta_vec + alpha_mean, sigma_resid)
credint2 = cbind(credint2, corrected)
colnames(credint2)[5] <- "corrected"
credint2
```

	2.5%	50%	97.5%	true_vals	corrected
## beta_mean	18.39846591	19.24987218	19.8150943	14.9048326	11.1025994
## alpha_mean	-7.62067516	0.48512666	9.0095442	-9.7067334	0.5393192
## eta_beta[1]	-0.67366840	0.02003857	0.9514246	2.2660105	9.4727995
## eta_beta[2]	-1.15214690	-0.05167200	0.8601764	9.4727995	11.4673746
## eta_beta[3]	-2.05418282	-0.16017574	0.5474489	-3.4374580	24.1947858
## eta_beta[4]	-1.69526090	-0.06030707	0.8778490	9.2899531	12.9210723
## eta_beta[5]	-1.03501377	-0.02993664	0.9084060	-1.9837603	-0.6105675
## eta_beta[6]	-0.77816191	-0.02397415	0.8192952	-15.5154002	6.0191753
## eta_beta[7]	-0.35750953	0.22614854	1.6054094	-8.8856573	9.3427149
## eta_beta[8]	-0.36445860	0.20230848	1.4198086	-5.5621177	14.3337057
## eta_alpha[1]	-11.29123067	-2.78622721	5.3183983	-0.5711269	11.1525341
## eta_alpha[2]	-20.62970867	-12.03387012	-3.9357379	-3.7522985	-2.6197595
## eta_alpha[3]	-17.05161298	-8.53412948	-0.5660163	7.0869739	-10.8146268
## eta_alpha[4]	1.79700773	10.08681020	18.1409364	-1.1078934	-8.2606383
## eta_alpha[5]	-15.06070468	-6.40308856	1.6719936	1.4460951	11.0384900
## eta_alpha[6]	6.92352381	15.50354839	23.6114134	20.7452234	-6.3250744
## eta_alpha[7]	3.14886932	11.69229861	19.7627941	3.3816590	16.0285114
## eta_alpha[8]	-15.26917912	-6.77913182	1.2036460	25.7352448	11.9050031
## beta_sigma	0.03351243	0.40973732	1.5694478	21.6117366	-6.6373518
## alpha_sigma	6.69277102	10.43207398	19.7490321	3.0693816	0.9353610
## sigma_resid	0.69040320	0.80683548	0.9595048	0.9353610	11.1025994

we can see that few parameters were adequately recovered. This included alpha_mean and all of the eta_alphas (9 of 12 params)

5. Calculate the LOO-IC for each of your two models. Comment on which model is more parsimonious, given the data at hand (5 points).

Calculate the LOO and WAIC for a single model:

Full model

```
log_lik_1 = extract_log_lik(model_fit_1)
loo_1 = loo(log_lik_1)
loo_1$estimates
```

	Estimate	SE
## elpd_loo	-100.136957	6.099760
## p_loo	9.529479	1.676108
## looic	200.273913	12.199520

reduced, intercept RE model

```
log_lik_2 = extract_log_lik(model_fit_2)
loo_2 = loo(log_lik_2)
loo_2$estimates
```

	Estimate	SE
## elpd_loo	-100.7607	6.127462
## p_loo	11.8975	1.947453
## looic	201.5213	12.254924

model comparison of loo

```
loo::compare(loo_1, loo_2)
```

```
## elpd_diff      se
##      -0.6      0.9
```

Based on the p_{loo} values for both models, the estimate is 9.5 which is lower for the reduced, intercept only random effects model. Additionally, the $elpd_{loo}$ is lower but not significant because SE ranges overlap for both model estimates of $elpd_{loo}$.

Therefore, we can say the reduced model is more parsimonious but because p_{loo} is still greater than the number of parameters, there is a level of model misspecification remaining (other model architectures may improve performance).

Task 2 (15 points)


1. Subset your data to include only one of the groups.

```
# choosing Group 1 as the single group for this task
data_group1 = data_df[data_df$group_idx == 1, ]
```

2. For that specific group, create a scatterplot.

(For the following, use your most parsimonious model)

```
ggplot(data_group1, aes(x = x_scaled, y = y_obs)) +
  geom_point() +
  labs(title = 'Group 1',
       x = expression(X[scaled]),
       y = expression(Y)) +
  geom_smooth(method = "lm", se = FALSE, size = 0.2)
```



3. On the scatterplot, overlay the median model-fit (i.e., the median estimated line).


```
group_idx_rd = sample(c(1:n_group), 1)
data_df_sub = data_df[group_idx == group_idx_rd,]
median_set = summary(model_fit_1)$summary[1:12, '50%']

beta_median = median_set[1]
alpha_mean_median = median_set[2]
eta_alpha_media = median_set[2 + group_idx_rd]
```

4. Bootstrap from the posterior to overlay many more possible model-fits. Be sure that for each of these model-fits, you are drawing parameter sets from your MCMC samples.: To get a joint draw, first specify which sample you will take from your MCMC chains. For instance:

```
plot(data_df[,2],data_df[,1],pch=16,xlab="X",ylab="y")
for (i in 1:50) {
  temp_idx = sample(c(1:n_iter_total-n_burn), 1)
  temp_alpha_mean = model_out_1$alpha_mean[temp_idx]
  temp_eta_alpha = model_out_1$eta_alpha[temp_idx,group_idx_rd]
  temp_beta = model_out_1$beta[temp_idx]
  abline(a = temp_eta_alpha + temp_alpha_mean, b = temp_beta, col = scales::a
lpha("gray",alpha=0.5))
}

abline(a = alpha_mean_median + eta_alpha_media, b= beta_median,col="black")
```



5. Make sure it is easy to visually distinguish between the median model-fit and the other possible model-fits. (You might want to plot the "other" fits before overlaying the median fit).